

# 1 Converting array into heap

## Problem Introduction

In this problem you will convert an array of integers into a heap. This is the crucial step of the sorting algorithm called HeapSort. It has guaranteed worst-case running time of  $O(n \log n)$  as opposed to QuickSort's average running time of  $O(n \log n)$ . QuickSort is usually used in practice, because typically it is faster, but HeapSort is used for external sort when you need to sort huge files that don't fit into memory of your computer.

## Problem Description

**Task.** The first step of the HeapSort algorithm is to create a heap from the array you want to sort. By the way, did you know that algorithms based on Heaps are widely used for external sort, when you need to sort huge files that don't fit into memory of a computer?

Your task is to implement this first step and convert a given array of integers into a heap. You will do that by applying a certain number of swaps to the array. Swap is an operation which exchanges elements  $a_i$  and  $a_j$  of the array  $a$  for some  $i$  and  $j$ . You will need to convert the array into a heap using only  $O(n)$  swaps, as was described in the lectures. Note that you will need to use a min-heap instead of a max-heap in this problem.

**Input Format.** The first line of the input contains a single integer  $n$ . The next line contains  $n$  space-separated integers  $a_i$ .

**Constraints.**  $1 \leq n \leq 100\,000$ ;  $0 \leq i, j \leq n - 1$ ;  $0 \leq a_0, a_1, \dots, a_{n-1} \leq 10^9$ . All  $a_i$  are distinct.

**Output Format.** The first line of the output should contain single integer  $m$  – the total number of swaps.  $m$  **must satisfy conditions**  $0 \leq m \leq 4n$ . The next  $m$  lines should contain the swap operations used to convert the array  $a$  into a heap. Each swap is described by a pair of integers  $i, j$  – the 0-based indices of the elements to be swapped. After applying all the swaps in the specified order the array must become a heap, that is, for each  $i$  where  $0 \leq i \leq n - 1$  the following conditions must be true:

1. If  $2i + 1 \leq n - 1$ , then  $a_i < a_{2i+1}$ .
2. If  $2i + 2 \leq n - 1$ , then  $a_i < a_{2i+2}$ .

Note that all the elements of the input array are distinct. Note that any sequence of swaps that has length at most  $4n$  and after which your initial array becomes a correct heap will be graded as correct.

### Sample 1.

Input:

```
5
5 4 3 2 1
```

Output:

```
3
1 4
0 1
1 3
```

After swapping elements 4 in position 1 and 1 in position 4 the array becomes 5 1 3 2 4.

After swapping elements 5 in position 0 and 1 in position 1 the array becomes 1 5 3 2 4.

After swapping elements 5 in position 1 and 2 in position 3 the array becomes 1 2 3 5 4, which is already a heap, because  $a_0 = 1 < 2 = a_1$ ,  $a_0 = 1 < 3 = a_2$ ,  $a_1 = 2 < 5 = a_3$ ,  $a_1 = 2 < 4 = a_4$ .

**Sample 2.**

Input:

```
5
1 2 3 4 5
```

Output:

```
0
```

The input array is already a heap, because it is sorted in increasing order.

**Starter Files**

There are starter solution only for C++, Java and Python3, and if you use other languages, you need to implement solution from scratch. Starter solutions read the array from the input, use a quadratic time algorithm to convert it to a heap and use  $\Theta(n^2)$  swaps to do that, then write the output. You need to replace the  $\Theta(n^2)$  implementation with an  $O(n)$  implementation using no more than  $4n$  swaps to convert the array into heap.

**What To Do**

Change the BuildHeap algorithm from the lecture to account for min-heap instead of max-heap and for 0-based indexing.

## 2 Parallel processing

### Problem Introduction

In this problem you will simulate a program that processes a list of jobs in parallel. Operating systems such as Linux, MacOS or Windows all have special programs in them called schedulers which do exactly this with the programs on your computer.

### Problem Description

**Task.** You have a program which is parallelized and uses  $n$  independent threads to process the given list of  $m$  jobs. Threads take jobs in the order they are given in the input. If there is a free thread, it immediately takes the next job from the list. If a thread has started processing a job, it doesn't interrupt or stop until it finishes processing the job. If several threads try to take jobs from the list simultaneously, the thread with smaller index takes the job. For each job you know exactly how long it will take any thread to process it, and this time is the same for all the threads. You need to determine for each job which thread will process it and when it will start processing.

**Input Format.** The first line of the input contains integers  $n$  and  $m$ . The second line contains  $m$  integers  $t_i$  – the times in seconds it takes any thread to process the  $i$ -th job. The times are given in the same order as they are in the list from which threads take jobs. Threads are indexed starting from 0.

**Constraints.**  $1 \leq n \leq 10^5$ ;  $1 \leq m \leq 10^5$ ;  $0 \leq t_i \leq 10^9$ .

**Output Format.** Output exactly  $m$  lines.  $i$ -th line (0-based index is used) should contain two space-separated integers – the 0-based index of the thread which will process the  $i$ -th job and the time in seconds when it will start processing that job.

#### Sample 1.

Input:

```
2 5
1 2 3 4 5
```

Output:

```
0 0
1 0
0 1
1 2
0 4
```

1. The two threads try to simultaneously take jobs from the list, so thread with index 0 actually takes the first job and starts working on it at the moment 0.
2. The thread with index 1 takes the second job and starts working on it also at the moment 0.
3. After 1 second, thread 0 is done with the first job and takes the third job from the list, and starts processing it immediately at time 1.
4. One second later, thread 1 is done with the second job and takes the fourth job from the list, and starts processing it immediately at time 2.
5. Finally, after 2 more seconds, thread 0 is done with the third job and takes the fifth job from the list, and starts processing it immediately at time 4.

## Sample 2.

Input:

```
4 20
1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Output:

```
0 0
1 0
2 0
3 0
0 1
1 1
2 1
3 1
0 2
1 2
2 2
3 2
0 3
1 3
2 3
3 3
0 4
1 4
2 4
3 4
```

Explanation: Jobs are taken by 4 threads in packs of 4, processed in 1 second, and then the next pack comes. This happens 5 times starting at moments 0, 1, 2, 3, and 4. After that all the  $5 \times 4 = 20$  jobs are processed.

## Starter Files

The starter solutions for C++, Java and Python3 in this problem read the input, apply an  $\Theta(n^2)$  algorithm to solve the problem and write the output. You need to replace the  $\Theta(n^2)$  algorithm with a faster one. If you use other languages, you need to implement the solution from scratch.

## What To Do

Think about the sequence of events when one of the threads becomes free (at the start and later after completing some job). How to apply priority queue to simulate processing of these events in the required order? Remember to consider the case when several threads become free simultaneously.

Beware of integer overflow in this problem: use type `long` in C++ and type `long` in Java wherever the regular type `int` can overflow given the restrictions in the problem statement.

### 3 Merging tables

#### Problem Introduction

In this problem, your goal is to simulate a sequence of merge operations with tables in a database.

#### Problem Description

**Task.** There are  $n$  tables stored in some database. The tables are numbered from 1 to  $n$ . All tables share the same set of columns. Each table contains either several rows with real data or a [symbolic link](#) to another table. Initially, all tables contain data, and  $i$ -th table has  $r_i$  rows. You need to perform  $m$  of the following operations:

1. Consider table number  $destination_i$ . Traverse the path of symbolic links to get the data. That is, while  $destination_i$  contains a symbolic link instead of real data do
$$destination_i \leftarrow \text{symlink}(destination_i)$$
2. Consider the table number  $source_i$  and traverse the path of symbolic links from it in the same manner as for  $destination_i$ .
3. Now,  $destination_i$  and  $source_i$  are the numbers of two tables with real data. If  $destination_i \neq source_i$ , copy all the rows from table  $source_i$  to table  $destination_i$ , then clear the table  $source_i$  and instead of real data put a symbolic link to  $destination_i$  into it.
4. Print the maximum size among all  $n$  tables (recall that size is the number of rows in the table). If the table contains only a symbolic link, its size is considered to be 0.

See examples and explanations for further clarifications.

**Input Format.** The first line of the input contains two integers  $n$  and  $m$  – the number of tables in the database and the number of merge queries to perform, respectively. The second line of the input contains  $n$  integers  $r_i$  – the number of rows in the  $i$ -th table. Then following  $m$  lines describing merge queries. Each of them contains two integers  $destination_i$  and  $source_i$  – the numbers of the tables to merge.

**Constraints.**  $1 \leq n, m \leq 100\,000$ ;  $0 \leq r_i \leq 10\,000$ ;  $1 \leq destination_i, source_i \leq n$ .

**Output Format.** For each query print a line containing a single integer – the maximum of the sizes of all tables (in terms of the number of rows) after the corresponding operation.

#### Sample 1.

Input:

```
5 5
1 1 1 1 1
3 5
2 4
1 4
5 4
5 3
```

Output:

```
2
2
3
5
5
```

In this sample, all the tables initially have exactly 1 row of data. Consider the merging operations:

1. All the data from the table 5 is copied to table number 3. Table 5 now contains only a symbolic link to table 3, while table 3 has 2 rows. 2 becomes the new maximum size.
2. 2 and 4 are merged in the same way as 3 and 5.
3. We are trying to merge 1 and 4, but 4 has a symbolic link pointing to 2, so we actually copy all the data from the table number 2 to the table number 1, clear the table number 2 and put a symbolic link to the table number 1 in it. Table 1 now has 3 rows of data, and 3 becomes the new maximum size.
4. Traversing the path of symbolic links from 4 we have  $4 \rightarrow 2 \rightarrow 1$ , and the path from 5 is  $5 \rightarrow 3$ . So we are actually merging tables 3 and 1. We copy all the rows from the table number 1 into the table number 3, and now the table number 3 has 5 rows of data, which is the new maximum.
5. All tables now directly or indirectly point to table 3, so all other merges won't change anything.

### Sample 2.

Input:

```
6 4
10 0 5 0 3 3
6 6
6 5
5 4
4 3
```

Output:

```
10
10
10
11
```

Explanation: In this example, tables have different sizes. Let us consider the operations:

1. Merging the table number 6 with itself doesn't change anything, and the maximum size is 10 (table number 1).
2. After merging the table number 5 into the table number 6, the table number 5 is cleared and has size 0, while the table number 6 has size 6. Still, the maximum size is 10.
3. By merging the table number 4 into the table number 5, we actually merge the table number 4 into the table number 6 (table 5 now contains just a symbolic link to table 6) so the table number 4 is cleared and has size 0, while the table number 6 has size 6. Still, the maximum size is 10.
4. By merging the table number 3 into the table number 4, we actually merge the table number 3 into the table number 6 (table 4 now contains just a symbolic link to table 6), so the table number 3 is cleared and has size 0, while the table number 6 has size 11, which is the new maximum size.

### Starter Files

The starter solutions in C++, Java and Python3 read the description of tables and operations from the input, declare and partially implement disjoint set union, and write the output. You need to complete the implementation of disjoint set union for this problem. If you use other languages, you will have to implement the solution from scratch.

## What To Do

Think how to use disjoint set union with path compression and union by rank heuristics to solve this problem. In particular, you should separate in your thinking the data structure that performs union/find operations from the merges of tables. If you're asked to merge first table into second, but the rank of the second table is smaller than the rank of the first table, you can ignore the requested order while merging in the Disjoint Set Union data structure and join the node corresponding to the second table to the node corresponding to the first table instead in your Disjoint Set Union. However, you will need to store the number of the actual second table to which you were requested to merge the first table in the parent node of the corresponding Disjoint Set, and you will need an additional field in the nodes of Disjoint Set Union to store it.